

# Text Classification (Part III)

[DAT640] Information Retrieval and Text Mining

Krisztian Balog

University of Stavanger

August 27, 2019

# Recap

- Implementing a text classification model using `scikit-learn`
  - GitHub: `code/text_classification.ipynb`
- Word counts used as features
- Document-term matrix is huge, but most of the values are zeros; stored as a sparse matrix

	$t_1$	$t_2$	$t_3$	$\dots$	$t_m$
$d_1$	1	0	2		0
$d_2$	0	1	0		2
$d_3$	0	0	1		0
$\dots$					
$d_n$	0	1	0		0


# Discussion

## Question

What are possible shortcomings of using raw term frequencies?

# Zip's law

- Given some corpus of natural language utterances, the frequency of any word is inversely proportional to its rank in the frequency table
  - Word number  $n$  has a frequency proportional to  $1/n$



# English language

- Most frequent words
  - the (7%)
  - of (3.5%)
  - and (2.8%)
- Top 135 most frequent words account for half of the words used

# Term weighting

- Intuition #1: terms that appear often in a document should get high weights
  - E.g., The more often a document contains the term “dog,” the more likely that the document is “about” dogs
- Intuition #2: terms that appear in many documents should get low weights
  - E.g., stopwords, like “a,” “the,” “this,” etc.
- How do we capture this mathematically?
  - Term frequency
  - Inverse document frequency

# Term frequency (TF)

- We write  $c_{t,d}$  for the raw count of a term in a document
- **Term frequency**  $tf_{t,d}$  reflects the importance of a term ( $t$ ) in a document ( $d$ )
- Variants
  - Binary:  $tf_{t,d} \in \{0, 1\}$
  - Raw count:  $tf_{t,d} = c_{t,d}$
  - **L1-normalized**:  $tf_{t,d} = \frac{c_{t,d}}{|d|}$ 
    - where  $|d|$  is the length of the document, i.e., the sum of all term counts in  $d$ :  
 $|d| = \sum_{t \in d} c_{t,d}$
  - L2-normalized:  $tf_{t,d} = \frac{c_{t,d}}{\|d\|}$ 
    - where  $\|d\| = \sqrt{\sum_{t \in d} (c_{t,d})^2}$
  - Log-normalized:  $tf_{t,d} = 1 + \log c_{t,d}$
  - ...
- By default, when we refer to TF we will mean the L1-normalized version

# Inverse document frequency (IDF)

- **Inverse document frequency**  $idf_t$  reflects the importance of a term ( $t$ ) in a collection of documents
  - The more documents that a term occurs in, the less discriminating the term is between documents, consequently, the less “useful”

$$idf_t = \log \frac{N + 1}{n_t}$$

- where  $N$  is the total number of documents in the collection and  $n_t$  is the number of documents that contain  $t$
- Log is used to “dampen” the effect of IDF



# Term weighting (TF-IDF)

- Combine TF and IDF weights by multiplying them:

$$tfidf_{t,d} = tf_{t,d} \cdot idf_t$$

- Term frequency weight measures importance in document
- Inverse document frequency measures importance in collection

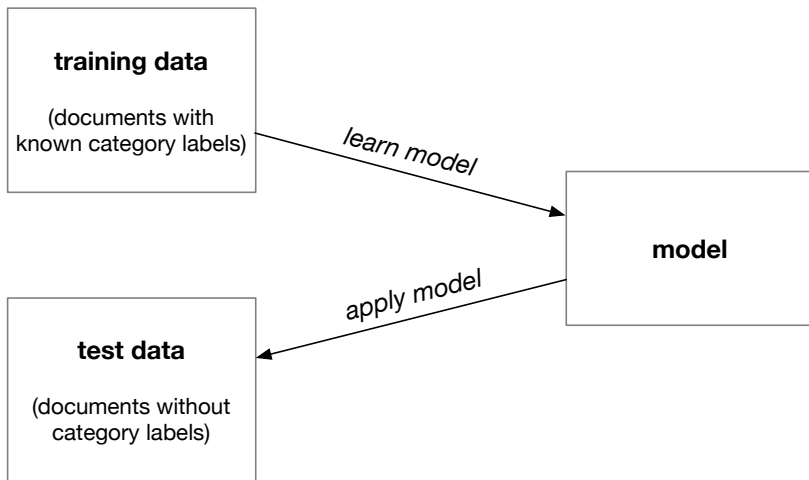
## Exercise #1 (paper-based)

**Create document-term matrix using TF-IDF weighting from a set of documents.**

# Code

- GitHub: `code/text_feature_extraction.ipynb`

# Text classification



# Text classification

- Formally: Given a training sample of documents  $\mathbf{X}$  and corresponding labels  $y$ ,  $((\mathbf{X}, y) = \{(x_1, y_1), \dots (x_n, y_n)\})$ , build a model  $f$  that can predict the class  $y' = f(x)$  for an unseen document  $x$
- Two popular classification models:
  - Naive Bayes
  - SVM

## Exercise #2 (coding)

- Compare two machine learning models and different term weighting schemes
  - Naive Bayes and SVM
  - Raw term count, TF weighting, and TF-IDF weighting
- Complete the TODOs and fill out the results table  
GitHub: `exercises/lecture_04/exercise_2.ipynb` (make a local copy)

# Assignment 1B